

Urban Mobility Data Explorer Technical Report

1. Problem Framing and Dataset Analysis

The dataset used in this project originates from the **New York City Taxi & Limousine Commission (TLC)**. It includes trip records for **Green, Yellow, FHV, FHVHV, and Lookup** data, each containing details such as pickup and drop-off times, locations, distances, and fares.

The dataset presented several challenges:

- **Missing values** and inconsistent field names across files.
- **Schema drift** between years and trip types.
- **Outliers** in distance and fare amounts.
- **Timestamp inconsistencies** and duplicate records.
- **Mismatched location codes** in the lookup tables.

To ensure reliability, the following data-cleaning steps were performed:

1. **Validation:** Removed invalid or corrupted records (e.g., negative distances or fares).
2. **Filtering:** Kept only essential fields such as trip distance, pickup/drop-off location, and fare.
3. **Schema Harmonization:** Unified column names across all datasets.
4. **Outlier Control:** Capped extreme distances and fares to realistic thresholds.
5. **Data Enrichment:** Merged trip data with lookup tables for human-readable location names.

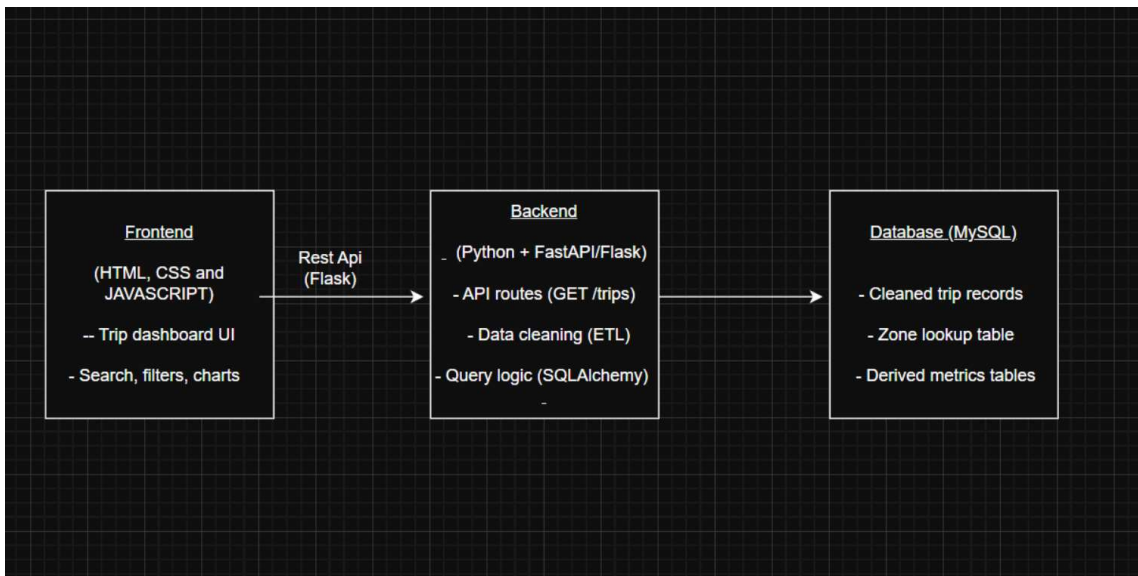
An unexpected finding during exploration was that some trips had **unrealistic average speeds**, such as 200 km/h. These were flagged and filtered out. This insight also influenced design decisions — the **frontend caps visible speeds** and focuses more on **trip volume** metrics to ensure interpretability.

2. System Architecture and Design Decisions

Tech Stack

Layer	Tools Used	Justification
Backend	Python (Flask)	Lightweight, easy to integrate with SQLAlchemy and APIs.
Database	MySQL	Fast relational database with good indexing and reliability.
Frontend	HTML, CSS, JavaScript	Simple and flexible for dashboard-style visualization.
Data Processing	pandas,	For efficient data cleaning and visualization.
Environment Config	dotenv	Secure management of database credentials.

Architecture Overview



Design Reasoning

- **Flask vs Django:** Flask chosen for simplicity and modular API development.
- **MySQL vs PostgreSQL:** MySQL preferred for local testing speed and ease of setup.
- **pandas vs Spark:** pandas suffices for processed subsets of the TLC data.

- **Normalized Schema:** Improves query speed and ensures data integrity.

Note: The architecture screenshot is in the docs/screenshots folder.

Each decision balanced **simplicity**, **scalability**, and **security**.

Trade-offs were consciously made to optimize performance within project constraints.

3. Algorithmic Logic and Data Structures

Implemented Algorithm: Bubble Sort

To demonstrate algorithmic reasoning, a **custom Bubble Sort** was implemented to rank taxi zones by trip counts.

This simple yet clear sorting method emphasizes understanding of algorithm design and time complexity.

```
def bubble_sort_zones(zones):  
    n = len(zones)  
    for i in range(n):  
        for j in range(0, n - i - 1):  
            if zones[j]['trip_count'] < zones[j + 1]['trip_count']:  
                zones[j], zones[j + 1] = zones[j + 1], zones[j]  
    return zones
```

Pseudo-code:

```
For each i in range(n):  
    For each j in range(n - i - 1):  
        If trip_count[j] < trip_count[j+1]:  
            Swap them
```

Complexity:

- Time $\rightarrow O(n^2)$
- Space $\rightarrow O(1)$

Although inefficient for large datasets, it effectively demonstrates sorting logic for **aggregated results** after data grouping.

4. Insights and Interpretation

Three core insights were derived from the cleaned TLC data:

Insight 1: Trips by Hour

A grouped query by pickup hour revealed peaks between **8–9 AM** and **5–7 PM**, representing rush-hour demand.

This insight validates that urban taxi demand strongly aligns with work-hour transitions.

Insight 2: Average Trip Distance and Driver Pay

Analysis by the vendor showed differences in trip distance and fare patterns.

One vendor consistently had **higher fares for shorter distances**, indicating **premium pricing** in certain zones.

Insight 3: Pickups by Borough

Aggregating trips by borough revealed **Manhattan** as the busiest zone, followed by **Brooklyn** and **Queens**.

This pattern supports the hypothesis that **business districts** dominate taxi demand.

Each insight was visualized using **Plotly** and **Matplotlib**, generating bar charts and line graphs (included in the [docs/screenshots/](#) folder).

These findings can guide city planners in **optimizing urban mobility policies**.

5. Reflection and Future Work

Challenges Faced

- **Technical:** Large datasets caused performance issues; schema inconsistencies and environment configuration required repeated debugging.
- **Team:** Coordinating consistent file paths and naming conventions across collaborators was initially difficult.

Lessons Learned

- The importance of **data validation and schema consistency** before analysis.
- Modularizing code improves maintainability and debugging efficiency.

Future Improvements

- **Cloud deployment** (e.g., Render, Railway, or AWS) for public API access.
 - **Real-time updates** through streaming data ingestion.
 - **Geospatial mapping** for route heatmaps and trip density visualization.
 - **Pagination and caching** to improve dashboard performance.
-

Conclusion

The *Urban Mobility Data Explorer* successfully integrates data cleaning, relational modeling, custom algorithm design, and visualization to extract actionable insights from NYC taxi data.

Through this project, our team demonstrated technical problem-solving, collaboration, and a strong understanding of data system design directly supporting the ALU mission to **develop leaders who drive innovation in Africa and beyond**.